

Predictive Information in Conjunction Data-Message Histories Decays as Time to Closest Approach Decreases

Arjun Vijay Prakash

City Montessori School, Kanpur Road Campus
Lucknow, India

Abstract—Conjunction Data Messages recompute collision probability as later tracking shrinks orbital uncertainty. An earlier forecast is useful only if the history available then contains information beyond copying the latest report. This study freezes the clock at 72, 48, 24, and 12 hours before closest approach on the public ESA Collision Avoidance Challenge archive and asks that question under a leakage-safe split. Extra information is real, but it is a floor-collapse statistic and a long-horizon statistic, not a typical-event accuracy win and not a high-risk warning win. On 1659 held-out events, 1352 later reports sit at the dataset floor of -30 , and persistence already has median absolute error 0. Unguarded XGBoost lowers mean absolute error from 5.080 to 2.809, yet floor-excluded MAE rises from 4.073 to 7.562 and a Wilcoxon test on paired absolute error does not reject equality ($p = 0.91$). The learned advantage is 4.53 log units at 72 hours and 0.06 at 12 hours. Validation therefore selects a two-part floor hurdle (test MAE 2.109, $F_2 = 0$). On ESA’s 2167 official-test events, clipped persistence matches Uriot last-risk-prediction loss $L = 0.694$; the floor hurdle lowers MAE but raises L to 70.5. Combined covariance volume tracks later movement; the max-risk gap tracks floor membership, not extra movement. This is a research prototype, not flight software.

Index Terms—conjunction data message, collision probability, space debris, forecast horizon, persistence

I. INTRODUCTION

Two orbiting objects can pass close enough to generate a conjunction alert. The reported collision probability P_c is not a physical constant: it is recomputed as radar and optical observations shrink (or fail to shrink) the combined covariance [1], [2]. Waiting usually improves the estimate and leaves less time to plan. ESA posed the forecasting form of this problem in the 2019 Collision Avoidance Challenge: predict the final reported $\log_{10} P_c$ from messages with time to closest approach of at least two days [3], [4]. Persistence—copying the latest pre-cutoff risk—was a strong baseline. Only 12 of 96 teams beat it on ESA-style loss.

That competition does not answer how the extra information, if any, depends on remaining time, nor whether a mean-error win is a typical-event win. Physics-based P_c forecasting via expected covariance reduction [5] and operational work on probability dilution [2], [6] suggest that large early covariances can make P_c look safely small, then later updates either collapse risk to negligible or concentrate it. A recent machine-learning model predicts which conjunctions will keep high covariance, for extra tracking [7]. The target here is different:

the later reported $\log_{10} P_c$ itself, under a frozen information cutoff.

The working hypotheses are as follows.

- H1: Extra predictive information beyond copying the latest report is largest at long horizons and approaches zero near closest approach.
- H2: Most of the mean-error win is rare floor jumps, not typical events becoming more accurate.
- H3: Extra information does not automatically improve high-risk decisions under the ESA class $\log_{10} P_c \geq -6$.
- H4: Large combined covariance at T-48 predicts later $|\Delta\text{risk}|$. The snapshot gap $\text{max_risk_estimate} - \text{risk}$ predicts floor membership, not larger later movement.

II. METHODS

A. Data

Labeled CDMs from the ESA training archive comprise 162,634 rows and 13,154 events, 2015–2019, anonymized [4]. Events need a message at or before the cutoff and a later labeled update. That yields 8,293 eligible events at 48 hours before closest approach. Train, validation, calibration, and test are event-disjoint (3,731 / 1,659 / 1,244 / 1,659). Seed 42 is the reported local split. Five grouped redraws use seeds 42–46 with the same recipe and no test-tuned knobs; those rows are extra measurements, not extra models. Leave-one-high-risk-out retrains residual XGBoost on all eligible events except one of the 66 local $\log_{10} P_c \geq -6$ events, then scores that event. Features use only rows with time to closest approach at or above the cutoff. The later $\log_{10} P_c$ is the label, never an input. Official-test inputs are the public Kelvins `test_data.csv` (all `time_to_tca` ≥ 2). Labels are the released `true_risk` column from Zenodo record 4463683 (25 January 2021), joined after freeze. Shared numeric `event_ids` between train and test are independent numbering, not shared conjunctions (0 identical pre-cutoff snapshots). Private post-cutoff fields are not features. No hyperparameter was changed after seeing official-test scores.

B. Model

The current selected policy on the project surface is a two-part floor hurdle: a classifier for $P(y = -30)$ from the final reported risk, a residual XGBoost fit on non-floor

training events only, threshold 0.15 chosen on validation, and no -6 persistence guard. At inference the model predicts -30 when $P(\text{floor})$ exceeds that threshold, otherwise risk + $\hat{\Delta}$. Live intervals are split conformal around that point, fit on the calibration split. SHAP explains the residual regressor. Baselines are persistence, the training-set median, Ridge, unguarded XGBoost, and the 18 August bootstrap ensemble with a -6 persistence guard (an exhibit snapshot, not a live mode). A residual candidate that predicts $y - \text{risk}$ on all cutoff-safe training rows remains a research artifact. Official-test scoring uses those frozen artifacts only. Predictions below -6 are clipped to -6.001 for ESA-style loss, matching Uriot *et al.* [3]. A class-threshold sweep reuses those frozen point forecasts and changes only the class cut at $\log_{10} P_c \in \{-8, -7, -6, -5, -4\}$; models are not retuned. For each cut, predictions below t are clipped to $t - 0.001$.

C. Metrics

Reported scores are mean and median absolute error in $\log_{10} P_c$, the share within 0.5 log units, and ESA-style loss $L = \text{MSE}_{HR}/F_2$ with $\beta = 2$ on the -6 class [3]. Operational LEO reaction is nearer -4 to -5 ; ESA scored -6 to keep enough positives. The sweep reports L , F_2 , $n+$, and a false-reassurance analogue (accepted forecast under the existing -6 abstention mask with $\hat{y} < t$ while $y \geq t$). L is omitted when $F_2 = 0$. An event is a floor event if the later report satisfies $y \leq -30 + 10^{-6}$. Floor-excluded MAE is MAE on the complement. Residual MAE is $\text{mean}|y - \text{risk}|$ versus $\text{mean}|\hat{y} - \text{risk}|$: how far the later report moved, versus how far the model moved the snapshot. A 95% interval on MAE(persistence)–MAE(model) uses 1000 event-resamples. A two-sided Wilcoxon signed-rank test compares per-event absolute errors. Split conformal uses absolute residuals $|y - \hat{y}|$ of the live floor-hurdle point on the 1244 calibration events [8]. Research tables also report conformal bands around the 18 August ensemble median. Official-test L is the secondary evaluation on 2167 events with 150 high-risk positives. The dilution probe is a logistic / Spearman baseline, not an extra XGBoost model: `dilution_gap = max_risk_estimate - risk` is excluded from the exhibit feature list. Spearman correlations are reported on the frozen test. A logistic model of floor collapse on `dilution_gap`, miss distance, and message count is fit on train only; AUC is on test. Quartile edges for the gap table come from train. This probe follows Alfano and CARA dilution [2], [6]; it does not predict residual covariance for extra tracking [7]. Publication figures in `docs/figures/` are regenerated by `plots.py` with SI-style axis labels.

Honest metrics, residual and floor models, split-conformal coverage, and the class-threshold sweep were scored on the frozen 18 August split. Validation selected the floor hurdle as the live CDM policy. The homepage shows that forecast with a conformal interval. It does not offer a threshold picker or a second “paper mode.” The laboratory page reports research tables, including the August ensemble snapshot.

TABLE I

FROZEN LOCAL TEST AT 48 HOURS ($n = 1659$; 1352 FLOOR EVENTS). MAE IN $\log_{10} P_c$. Δ IS MAE(PERSISTENCE)–MAE(MODEL) WITH A 95% BOOTSTRAP INTERVAL (1000 EVENT RESAMPLES). WILCOXON p IS TWO-SIDED ON PAIRED ABSOLUTE ERROR. SELECTED POLICY: FLOOR HURDLE (VALIDATION). GUARDED ENS.: BOOTSTRAP ENSEMBLE WITH A -6 PERSISTENCE GUARD.

System	MAE	Med.	Non-fl.	Floor	Δ [95%]	p
Persistence	5.080	0.000	4.073	5.308	—	—
Unguarded XGB	2.809	0.554	7.562	1.730	2.270 [1.930, 2.638]	0.91
Residual XGB	2.760	0.453	7.375	1.712	2.319 [1.977, 2.685]	0.92
Floor hurdle	2.109	0.000	9.311	0.474	2.971 [2.524, 3.415]	$< 10^{-33}$
Guarded ens.	3.059	0.473	7.332	2.088	2.021 [1.716, 2.380]	0.31

III. RESULTS

Table I reports the frozen local test with floor slices and intervals. Of 1659 events, 1352 later reports sit at -30 . Persistence median AE is 0: more than half of reports do not move after 48 hours. Reconstructing $\hat{y} = \text{risk} + y - \text{risk}$ tests whether anything besides persistence exists. Residual XGBoost reaches test MAE 2.760 ($\Delta = 2.319$ [1.977, 2.685]), close to unguarded level-valued XGBoost (2.809). The two-part floor model reaches test MAE 2.109 ($\Delta = 2.971$ [2.524, 3.415]) and median AE 0. Floor-class confusion at the validation threshold 0.15 is 1319 true positives, 174 false positives, 33 false negatives, and 133 true negatives (recall 0.976, precision 0.883). Floor-only MAE falls to 0.474, but floor-excluded MAE rises to 9.311 versus persistence 4.073: the mean and Wilcoxon wins ($p < 10^{-33}$) come from calling the floor, including on 174 non-floor events. On validation, this candidate ranks first by MAE (1.916 versus unguarded XGBoost 2.669) and is the live CDM policy. $F_2 = 0$ for the floor model, residual model, and unguarded XGBoost. ESA-style loss still ties at 0.167 for persistence and the 18 August ensemble snapshot (H3). Fig. 1 shows why the mean win is not a typical-event win: most later reports do not move (spike at $y - \text{risk} = 0$), while a left tail of floor collapses to -30 dominates MAE. Across five grouped redraws, unguarded XGBoost MAE advantage is 2.34 ± 0.04 and floor-excluded MAE is 7.78 ± 0.21 (persistence floor-excluded 4.32 ± 0.23). Residual reconstruction is 2.35 ± 0.04 MAE advantage. Seed 42 remains the reported row in Table I. Leave-one-out on the 66 high-risk eligible events: persistence is closer than residual on 66/66 (mean absolute error 1.21 versus 12.59). The constant training-set median still beats the Guarded ens. on overall MAE (3.002) but not the live floor hurdle (2.109).

Table II and Fig. 2 support H1 on this freeze: the learned advantage shrinks as closest approach approaches. Whiskers at T=48 are the 95% bootstrap interval on MAE(persistence)–MAE(XGBoost), drawn on the XGBoost point. Five grouped redraws give a 2.34 ± 0.04 advantage at 48 h; that redraw s.d. is smaller than the marker. The 12-hour gap is 0.06 log units and does not yet have a redraw interval.

Snapshot-only MAE is 2.904; adding historical summaries of the same fields lowers it to 2.851; covariance trends reach 2.808. History therefore adds a small increment that also still




Fig. 1. Frozen test histograms of later movement $y - \text{risk}$ and unguarded XGBoost residual $y - \hat{y}$. The spike at 0 is reports that do not move. The left tail is collapse to the dataset floor -30 .

TABLE II
SINGLE XGBOOST VERSUS PERSISTENCE BY HORIZON. MAE IN $\log_{10} P_c$. THE 72 H AND 12 H ELIGIBLE SETS OVERLAP THE 48 H SPLIT BUT ARE NOT IDENTICAL.

Horizon	XGBoost	Persist.	Adv.
72 h	3.214	7.748	4.534
48 h	2.808	5.080	2.272
24 h	2.110	2.634	0.524
12 h	1.384	1.444	0.060

needs an interval.

Live split-conformal 90% intervals around the floor-hurdle point cover 90.1% of the test (mean width 21.03; $q_{0.1} = 10.51$). The 50% conformal radius is 0 because most calibration residuals are exact floor hits (empirical 50% coverage 79.5%). Table III and Fig. 3 still record the August-snapshot comparison. On the live conformal abstention rule, test coverage is 90.4% (159 of 1659 sent to review) and false reassurance is 2 of 9. A four-mission hold-out has high-risk MAE 19.2 on a single held-out high-risk event.

Table IV and Fig. 4 report the one-shot official-test score. Persistence $L = 0.694$ matches the last-risk-prediction number in Uriot *et al.* on this distribution. The August ensemble snapshot ties that loss because the -6 guard copies the snapshot whenever it is already high-risk. Residual and floor candidates still reduce MAE (3.476 and 3.177 versus 5.209) and the floor hurdle Wilcoxon versus persistence is $p < 10^{-17}$, but F_2 collapses and L rises to 104 and 70. Unguarded XGBoost has $F_2 = 0$. None of these frozen models beats the published sesc $L = 0.556$. That is the result; the models were not returned.

Table V and Fig. 5 report H4. Combined logdet of the snapshot covariance correlates with later $|y - \text{risk}|$ ($\rho = 0.399$, $p < 10^{-63}$). F_{10} is weaker ($\rho = 0.108$). The max-risk gap is strongly associated with ending at the floor ($\rho = 0.407$;




Fig. 2. Forecast-horizon MAE for unguarded XGBoost and persistence. Whiskers at 48 h are the 95% bootstrap interval on the MAE gap, shown on the XGBoost point.

TABLE III
INTERVAL COVERAGE ON THE FROZEN LOCAL TEST ($n = 1659$). BOOTSTRAP BANDS ARE MODEL SPREAD. SPLIT CONFORMAL IS FIT ON 1244 CALIBRATION EVENTS ONLY.

Band	Nom.	Cover.	Width
Bootstrap spread	50%	26.0%	0.90
Bootstrap spread	90%	47.7%	2.05
Split conformal	50%	49.6%	0.93
Split conformal	90%	89.7%	18.33

logistic test AUC 0.819) but with *less* later movement ($\rho = -0.768$): large gaps are mostly snapshots already near -30 . Quartile 1 (small gap) has mean $|y - \text{risk}| = 13.03$ and floor rate 0.56; quartiles 3–4 have floor rates above 0.95 and mean movement below 0.75. The naive story that a large $\text{max_risk} - \text{risk}$ gap means the report will still move is rejected on this freeze. Large covariance volume predicting more movement is supported.

Table VI reuses the frozen local-test forecasts and changes only the class definition. Operational LEO reaction is nearer -4 to -5 ; ESA scored -6 so the positive class was large enough to estimate F_2 . Persistence and the August ensemble snapshot still tie at -6 ($n+ = 9$, $L = 0.167$, $F_2 = 0.361$, one accepted false reassurance) and at -5 ($n+ = 3$, $L = 0.032$, $F_2 = 0.349$, zero accepted false reassurances). At -8 and -7 persistence has higher F_2 and lower L than the ensemble. At -4 a single test event is positive and every frozen system misses it ($F_2 = 0$). Unguarded XGBoost, residual reconstruction, and the floor hurdle have $F_2 = 0$ at -7 through -4 .

IV. DISCUSSION

The extra information in a T-48 CDM history is not a general improvement in forecasting later $\log_{10} P_c$. It is early, and it is concentrated in later collapses to the dataset floor. That reading unifies H1 and H2 with the physical story already




Fig. 3. Nominal versus empirical coverage for bootstrap spread and split conformal (frozen test).

TABLE IV

OFFICIAL TEST ($n = 2167$; 150 HIGH-RISK; 1673 FLOOR). MAE IN $\log_{10} P_c$. $L = \text{MSE}_{HR}/F_2$ AFTER CLIPPING PREDICTIONS BELOW -6 TO -6.001 . FROZEN BEFORE LOOK.

System	MAE	Med.	Non-fl.	L	F_2
Persistence (LRP)	5.209	0.000	4.287	0.694	0.739
Unguarded XGB	3.504	0.690	9.333	1.75×10^6	0.000
Residual XGB	3.476	0.594	9.289	104	0.017
Floor hurdle	3.177	0.000	11.832	70.5	0.025
Guarded ens.	3.107	0.439	6.750	0.694	0.739

in the conjunction literature. Large early covariance can dilute P_c [6]. Later contraction often drops reported risk [2]. When the clock is stopped far from closest approach, a learned model can still tell some of those collapses from persistence. By 12 hours the latest report has already absorbed most of that update, and the learned advantage is 0.06 log units.

H3 follows from the same anatomy. ESA-style loss cares about the $\log_{10} P_c \geq -6$ tail. The floor hurdle, which validation selects because it wins on MAE, almost never predicts into that tail ($F_2 = 0$ locally; $L = 70.5$ on official test). The guarded ensemble ties persistence on L because it copies today's report once that report is already high-risk. Those are different jobs. Selecting on mean error ships a floor-calling policy. Selecting on ESA-style loss would have kept persistence, or the guard. The 2019 challenge already showed that beating persistence on that loss is hard [3]. This freeze does not overturn that. It says why a mean-error win can coexist with a worse high-risk score.

H4 is mixed in a way that is also consistent. Covariance volume at T-48 tracks later $|y - \text{risk}|$. The max-risk gap tracks whether the later report is already a floor event, not whether it will still move. That is a property of the CDM fields, not a




Fig. 4. Official-test ESA-style loss for frozen models versus Uriot last-risk-prediction (0.694) and sesc (0.556). Log scale.

TABLE V

FROZEN TEST BY TRAIN-QUARTILE OF
DILUTION_GAP = MAX_RISK_ESTIMATE - RISK ($n = 1659$).

Q	n	Floor	$ \Delta \text{risk} $	XGB MAE
1	410	0.556	13.03	7.35
2	416	0.781	6.15	2.90
3	417	0.964	0.52	0.69
4	414	0.954	0.75	0.37

covariance-tasking model [7].

The selected policy is therefore an honest floor-collapse forecast with a wide conformal band, not an operational warning system. Leave-one-out on the 66 local high-risk events still prefers persistence every time. This freeze uses historical ESA-supported missions from 2015 to 2019, not live catalogues. Manoeuvres are out of scope. The models were not retuned on official-test scores.

ACKNOWLEDGMENT

This work uses the ESA Space Debris Office Collision Avoidance Challenge dataset [3], [4]. ESA, ISRO, NASA, ISTRAC, and CCSDS did not endorse this project.

REFERENCES

- [1] Consultative Committee for Space Data Systems, *Conjunction Data Message (CCSDS 508.0-B-1)*, 2013.
- [2] M. D. Hejduk and L. C. Johnson, "Satellite conjunction assessment risk analysis for dilution region events: Issues and operational approaches," in *Space Traffic Management Conference*, 2016.
- [3] T. Uriot, D. Izzo, L. F. Simões, R. Abay, N. Einecke, S. Rebhan, J. Martinez-Heras, F. Letizia, J. Siminski, and K. Merz, "Spacecraft collision avoidance challenge: Design and results of a machine learning competition," *Astrodynamics*, vol. 6, pp. 121–140, 2022, arXiv:2008.03069.
- [4] ESA Advanced Concepts Team and Space Debris Office, "Collision avoidance challenge dataset," Zenodo, 2021.
- [5] M. Duncan, J. Wysack, and J. Frisbee, "Collision probability forecasting using a Monte Carlo simulation," in *Advanced Maui Optical and Space Surveillance Technologies Conference (AMOS)*, 2014.
- [6] S. Alfano, "Relating position uncertainty to maximum conjunction probability," in *AAS/AIAA Astrodynamics Specialist Conference*, 2003, AAS 03-548.




Fig. 5. Floor rate (bars, left axis) and mean $|y - \text{risk}|$ (line, right axis) by train-quartile of the max-risk gap (frozen test).

TABLE VI

FROZEN LOCAL TEST ($n = 1659$). SAME PREDICTIONS; CLASS CUT t ONLY. P IS PERSISTENCE; E IS THE AUGUST ENSEMBLE SNAPSHOT. FR IS ACCEPTED $\hat{y} < t$ WHILE $y \geq t$ UNDER THE AUGUST -6 ABSTENTION MASK. L OMITTED WHEN $F_2 = 0$.

t	n_+	$P F_2$	$P FR$	$P L$	$E F_2$	$E FR$	$E L$
-8	68	0.660	2	1.66	0.210	5	9.49
-7	37	0.502	3	0.777	0.215	4	2.51
-6	9	0.361	1	0.167	0.361	1	0.167
-5	3	0.349	0	0.032	0.349	0	0.032
-4	1	0	1	—	0	1	—

- [7] J. Stauch, T. Olson, C. Reid, C. Grey, D. C. Kesler, and B. G. Marchand, "Contextual predictive model for early identification of high-covariance conjunctions," *The Journal of the Astronautical Sciences*, vol. 73, p. 19, 2026.
- [8] J. Lei, M. G'Sell, A. Rinaldo, R. J. Tibshirani, and L. Wasserman, "Distribution-free predictive inference for regression," *Journal of the American Statistical Association*, vol. 113, no. 523, pp. 1094–1111, 2018.